

# CS 451 – Operating Systems

## Project 2

**Due: 11:59 PM on Tuesday, 25 August 2020**

### Objectives

- To familiarize yourself with the details of a MLFQ scheduler.
- To implement a basic MLFQ scheduler and FIFO scheduling method.
- To create system-call that extracts process states.

### Teaming up!

For this project, you have the option to work with a partner. Please email the instructor with your teammate's name and a team name (if you prefer one).

### Overview

In this project, you'll be implementing a simplified multi-level feedback queue (MLFQ) scheduler in xv6. The basic idea is simple.

- Build an MLFQ scheduler with **four priority queues**;
- the top queue (numbered 3) has the highest priority and the bottom queue (numbered 0) has the lowest priority.
- When a process uses up its time-slice (counted as a number of ticks), it should be downgraded to the next (lower) priority level.
- The time-slices for higher priorities will be shorter than lower priorities.
- The scheduling method in each of these queues will be round robin, except the bottom queue which will be implemented as FIFO.

### Useful Resources

[Chapter 5](#) in xv6 book.

### Project Details

You have two specific tasks for this part of the project. However, before starting these two tasks, you need first have a high-level understanding of how scheduler works in xv6.

Most of the code for the scheduler is quite localized and can be found in **proc.c**, where you should first look at the routine **scheduler()**. It's essentially looping forever and for each iteration, it looks for a runnable process across the ptable. If there are multiple runnable processes, it will select one according to some policy. As you observed in **lab 3 and lab 4**, the vanilla xv6 does no fancy

things about the scheduler; it simply schedules processes for each iteration in a round-robin fashion. For example, if there are three processes A, B and C, then the pattern under the vanilla round-robin scheduler will be A B C A B C ... , where each letter represents a process scheduled within **a timer tick**, which is essentially ~10ms, and you may assume that this **timer tick is equivalent to a single iteration of the for loop** in the `scheduler()` code.

**Why 10ms?** This is based on the timer interrupt frequency setup in xv6 and you may find the code for it in `timer.c`.

Now to implement **MLFQ**, you need to schedule the process for **some time-slice**, which is some multiple of timer ticks. For example, if a process is on the **highest priority level**, which has a time-slice of **8 timer ticks**, then you should schedule this process for **~80ms**, or equivalently, for **8 iterations**.

xv6 can perform a context-switch every time a timer interrupt occurs. For example, if there are 2 processes A and B that are running at the highest priority level (queue 3), and if the round-robin time slice for each process at level 3 (highest priority) is **8 timer ticks**, then if process A is chosen to be scheduled before B, A should run for a complete time slice (~80ms) before B can run. Note that even though process A runs for 8 timer ticks, every time a timer tick happens, process A will yield the CPU to the scheduler, and the scheduler will decide to run process A again (**until its time slice is complete**).

## Implement MLFQ

Your MLFQ scheduler must follow these very precise rules:

- Four priority levels numbered from 3 (highest) down to 0 (lowest).
- Whenever the xv6 10 ms timer tick occurs, the highest priority ready process is scheduled to run.
- The highest priority ready process is scheduled to run whenever the previously running process exits, sleeps, or otherwise yields the CPU.
- If there are more than one processes on the same priority level, then you scheduler should schedule all the processes at that particular level in a round robin fashion. Except for priority level 0, which will be scheduled using FIFO basis.
- When a timer tick occurs, whichever process was currently using the CPU should be considered to have used up an entire timer tick's worth of CPU, even if it did not start at the previous tick (Note that a timer tick is different than the time-slice.)

- The time-slice associated with priority 3 is 8 timer ticks; for priority 2 it is 16 timer ticks; for priority 1 it is 32 timer ticks, and for priority 0 it executes the process until completion.
- When a new process arrives, it should start at priority 3 (highest priority).
- At priorities 3, 2, and 1, after a process consumes its time-slice it should be downgraded one priority. At priority 0, the process should be executed to completion.
- If a process voluntarily relinquishes the CPU before its time-slice expires at a particular priority level, its time-slice should not be reset; the next time that process is scheduled, it will continue to use the remainder of its existing time-slice at that priority level.
- To overcome the problem of starvation, we will implement a mechanism for priority boost. If a process has waited 10x the time slice in its current priority level, it is raised to the next higher priority level at this time (unless it is already at priority level 3). For priority 0, which does not have a time slice, processes that have waited 500 ticks should be raised to priority 1.

## Create new system calls

You'll need to create one system call for this project: `int getpinfo(struct pstat *)` Because your MLFQ implementations are all in the kernel level, you need to extract useful information for each process by creating this system call so as to better test whether your implementation works as expected.

To be more specific, this system call returns 0 on success and -1 on failure. If success, some basic information about each process: its process ID, how many timer ticks have elapsed while running in each level, which queue it is currently placed on (3, 2, 1, or 0), and its current procstate (e.g., SLEEPING, RUNNABLE, or RUNNING) will be filled in the pstat structure as defined in the next page.

```
struct pstat {
    int inuse[NPROC]; // whether this slot of the process table is in use (1 or 0)
    int pid[NPROC];   // PID of each process
    int priority[NPROC]; // current priority level of each process (0-3)
    enum procstate state[NPROC]; // current state (e.g., SLEEPING or RUNNABLE) of each process
    int ticks[NPROC][4]; // number of ticks each process has accumulated at each of 4 priorities
    int wait_ticks[NPROC][4]; // number of ticks each process has waited before being scheduled
};
```

The file `pstat.h` is uploaded along with this document. Do not change the names of the fields in `pstat.h`

## Tips

Most of the code for the scheduler is quite localized and can be found in `proc.c`; the associated header file, `proc.h` is also quite useful to examine. To change the scheduler, not too much needs to be done; study its control flow and then try some small changes.

As part of the information that you track for each process, you will probably want to know its current priority level and the number of timer-ticks it has left.

It is much easier to deal with fixed-sized arrays in `xv6` than linked-lists. For simplicity, it's recommended that you use arrays to represent each priority level.

You'll need to understand how to fill in the structure `pstat` in the kernel and pass the results to user space and how to pass the arguments from user space to the kernel. Good examples of how to pass arguments into the kernel are found in existing system calls. In particular, follow the path of `fileread()`, which will lead you to `sys_read()`, which will show you how to use

`int argint(int n, int *ip)` in `syscall.c` (and related calls) to obtain a pointer that has been passed into the kernel. In addition to these files, refer to the **Code: System call** section in chapter 3 of the **xv6 documentation**.

## Testing

Testing is critical. Testing your code to make sure it works is crucial.

### Basic Test

You can test your MLFQ scheduler by writing workloads and instrumenting it with `getpinfo` system call. Attached is an example of a user program (`test.c`) which spins for a user input.

**`getpinfo()`** : The function returns some basic information about each running process, including how long it has run at each priority (measured in clock ticks) and its process ID and possibly the state as well.